

Lakshya Dhawan

(Ghaziabad, Uttar Pradesh)

Portfolio: vercel.app

LinkedIn: linkedin.com/in/lakshya19

Email: lakshyad1902@gmail.com

Mobile: +91-93104-91273

LeetCode: leetcode.com/lakshyad1902

SUMMARY

Junior Data Engineer with 1.5 years of experience in ETL pipeline monitoring and data platform support at KPMG India. Skilled in Python, SQL, and Azure Synapse, with hands-on experience in backend API development and data processing workflows. Recently built end-to-end data systems including a RAG-based pipeline integrating vector search and LLMs. Actively strengthening expertise in Databricks and distributed data processing. Seeking Data Engineering roles to build scalable data pipelines and analytics systems.

SKILLS

- **Languages:** Python, SQL
- **Data Engineering:** ETL Pipelines, Data Processing, Data Modeling (Basic), Batch Processing
- **Data & Cloud:** Azure Synapse, Databricks (Learning), AWS (Foundational), PostgreSQL, MongoDB
- **AI & ML:** RAG Systems, LLM Integration, Embeddings, FAISS, Transformers (HuggingFace)
- **Frameworks & Tools:** FastAPI, Pandas, PySpark (Basic), Docker, Git, Postman

EDUCATION

- **ABES Engineering College (AKTU)** Uttar Pradesh, India
B.Tech – Computer Science & Engineering (Data Science); GPA: 7.83 *Sept 2021 – May 2025*
- **Sanskar The Co-Educational School** Uttar Pradesh, India
Senior Secondary (CBSE); Percentage: 97.4% *2019 – 2021*

EXPERIENCE

- **KPMG India** Gurgaon, India
Analyst – Data Operations *Feb 2025 – Present*
 - **ETL Pipeline Monitoring (Azure Synapse):** Monitored and supported ETL pipelines in Azure Synapse, identifying and resolving data processing failures to ensure reliable data delivery to downstream systems.
 - **Data Debugging & Transformation Support:** Analyzed pipeline failures and debugged transformation logic in notebooks, improving data consistency and reducing recurring processing issues.
 - **Backend API Development (FastAPI):** Developed REST APIs for internal modules, integrating SQL-based data retrieval and enabling structured data access for application features.
 - **Cloud Migration Support (Azure → AWS):** Contributed to cloud migration planning by mapping Azure Synapse workflows to AWS services (S3, Glue, Lambda), supporting transition to scalable data architecture.
 - **GenAI Platform Support:** Validated document processing pipelines and API integrations for internal AI platform, ensuring correctness and reliability of data workflows.

PROJECTS

- **RAG-based Accountability Assistant (Python, FAISS, FastAPI, SQLite, Ollama):**
 - Designed a stateful Retrieval-Augmented Generation (RAG) system with separate ingestion and retrieval pipelines for context-aware user interactions.
 - Built ingestion pipeline to process user inputs via chunking, embedding (MiniLM), and indexing in FAISS, with metadata stored in SQLite for efficient mapping.
 - Implemented semantic retrieval using cosine similarity (FAISS) to fetch top-K relevant historical context for each query.
 - Integrated self-hosted LLM (Mistral via Ollama) to generate structured, grounded responses using prompt engineering and context injection.
 - Optimized system performance by persisting FAISS indices and minimizing redundant computation during retrieval.
 - Designed API endpoints using FastAPI for real-time query handling and response generation.
- **Reddit Sentiment Analysis Pipeline (FastAPI, Transformers, Python):**
 - Built a data processing pipeline to ingest Reddit posts using PRAW, perform text preprocessing, and classify sentiment using DistilBERT.
 - Implemented batch-style processing of external data with normalization and aggregation of sentiment scores.
 - Designed REST API using FastAPI to serve real-time analysis results with structured JSON outputs.
 - Applied text cleaning, tokenization, and input truncation (512 tokens) to optimize model performance.
 - Identified and addressed system limitations including lack of caching and synchronous model loading, proposing optimizations for scalability.

COURSES & CERTIFICATIONS

- **Databricks Data Engineer Associate** – In Progress (Expected April 2026)
- **AI-900** Azure AI Fundamentals by **Microsoft**
- Career Essentials in Generative AI by **Microsoft** and **LinkedIn**
- AWS Academy Cloud Foundations by **Amazon Web Services**
- SQL Intermediate Certificate by **HackerRank**
- Python for Data Science by **IBM**